

Kuan-Chen Chen

champion3.chen@gmail.com | Gainesville, FL | 352-246-3469

ckcinnabar.github.io | <https://www.linkedin.com/in/kuan-chen-chen/> | <https://github.com/Ckcinnabar/>

EDUCATION

University of Florida

Master of Science in Applied Data Science

Gainesville, FL

8/2024 - 5/2026

National Sun Yat-Sen University

Bachelor of Science in Applied Mathematics

Kaohsiung, Taiwan

9/2019 - 6/2023

Honors & Awards: Valedictorian, Department of Applied Mathematics; Distinguished Service Award

EXPERIENCE

NVIDIA NeMo Automodel ([MiniCPM5-1B](#), [Falcon H1](#))

Open-Source Contributor

5/2026-Present

- Contributed 2 merged PRs to NVIDIA's distributed LLM training library, adding 7 fine-tuning recipes (SFT, LoRA) for models from 0.5B to 34B
- Configured LoRA hyperparameters (rank, alpha, target modules) and tensor parallelism for 34B-scale training; validated convergence on T4 GPU runs

Florida Gators Women's Basketball, University of Florida

Gainesville, FL

Data Scientist Intern

8/2025 - 5/2026

- Replaced a manually-weighted player evaluation system with OA-Rating, an XGBoost + SHAP model validated against the legacy system ($r = 0.78$) across 120K+ observations, 5,500+ players, 350+ teams
- Built Elo-based team rating system with 83% agreement with NCAA's official national rankings, and clustered teams into 4 tactical archetypes to support real-time game preparation and opponent scouting
- Shipped production R Shiny dashboard that contributed to team improvement improvement from 60th nationally (16-17) to 48th nationally (18-15)

Bioinformatics Lab, University of Florida

Gainesville, FL

Student Research Assistant

12/2024 - 5/2026

- Designed stacking ensemble reaching 87.6% prediction accuracy (ROC-AUC 0.876) in distinguishing brain- vs. blood-derived sequences, outperforming the best individual model (85%) on 10,486 CDR3 sequences
- Developed biologically-informed data augmentation algorithm using BLOSUM62 substitution matrix, generating conservative sequence variants to expand minority class $10\times$ ($432 \rightarrow 4,622$) and mitigate severe class imbalance ($1:23 \rightarrow 1:1.5$)
- Engineered 42 physicochemical features from protein sequences and optimized decision thresholds under class imbalance, improving precision to 0.716 for high-confidence, interpretable predictions

PROJECTS

CV Generator (JD-Tailored Résumé Agent) [Ckcinnabar/CV Generator](#)

Gainesville, FL

- Built a hybrid retrieval pipeline — dense embeddings (Chroma) + BM25, fused via Reciprocal Rank Fusion — with a Claude Haiku selector ranking the top-5 most relevant items per JD
- Orchestrated via a LangGraph stateful graph with a human-in-the-loop interrupt for approval, and constrained Claude generation with hard-coded titles/dates to prevent hallucination
- Built an LLM supervisor that fact-checks every claim against source documents, plus a render-retry loop (docx2pdf $\rightarrow$ pypdf) that shortens content until it fits one page

cAIuldrón (Recipe Generation LLM) [Ckcinnabar/cAIuldrón](#)

Gainesville, FL

- Built an AI recipe app that turns ingredient photos into complete recipes with AI-generated images; shipped two production versions, cutting pipeline latency from 40-80s to 10-20s
- Developed v1 pipeline using CLIP and DETR for multi-ingredient detection across 525+ food classes (88% F1); fine-tuned Llama 3.2 1B via QLoRA on 7,913 curated RecipeNLG recipes, updating only 0.23% of parameters (2.8M) while achieving 92/100 recipe quality score
- Rebuilt v2 on LangGraph with Groq-hosted Llama 4 Scout (vision) and Llama 3.3 70B (generation); parallel fork/join ChromaDB RAG + Tavily search, removing the 4-6 GB VRAM dependency

SKILLS

Languages & ML: Python, SQL, R, Docker | PyTorch, scikit-learn, XGBoost, LightGBM, SHAP

LLM & GenAI: HuggingFace Transformers, LoRA/QLoRA, LangGraph, LangSmith, ChromaDB

Web & Viz: Flask, React.js, Gradio, Dash, Shiny, Power BI, Tableau, FastAPI | DB: MySQL, MongoDB